Titanic: Machine Learning from Disaster – Kaggle Competition

Gareth Austen

Data Science

General Assembly

April 2015

Agenda

1. Defining the Problem
2. Overview of Kaggle Data and Data Manipulation
3. Model Selection and Methodology
4. Results of the Modeling process
5. Future Steps!

Overview of the Problem

- One of the most infamous shipwrecks in history
- 1502 out of 2224 passengers were killed when the Titanic sunk
- Some groups are more likely to survive – eg women and children
- The Challenge: To predict passengers likelihood to survive

Overview of Kaggle Data and Data Manipulation

Overview of the Data

- Kaggle provides the data in the form of two csv's – one training and one test set
- Both datasets contain the following variables:
  - Survival
  - Pclass - Passenger Class
  - Name
  - Sex
  - Age
  - Sibsp
  - parch
  - Ticket
  - Fare
  - Cabin
  - Cabin Number
  - Port of Embarkation

Male vs Female Survival Rate

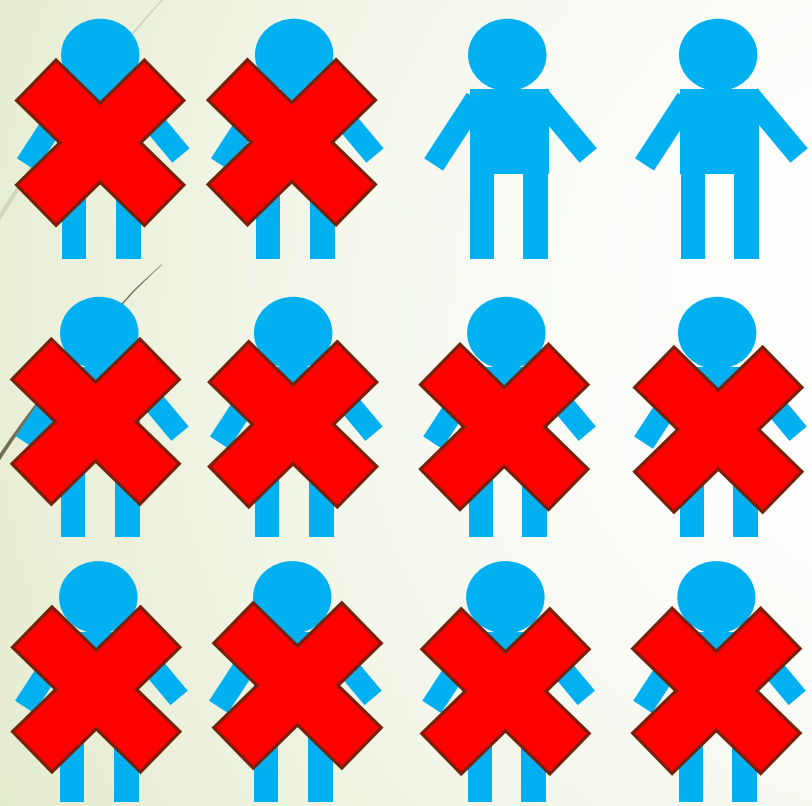

Male Survival Rate: 18%

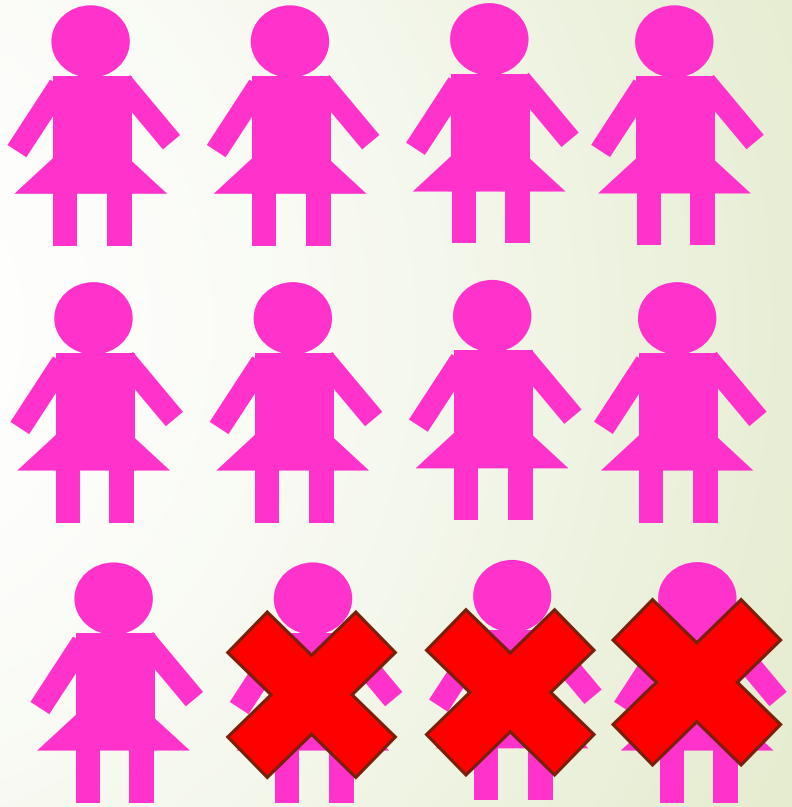

Female Survival Rate: 75%

Cross – Section of the Titanic

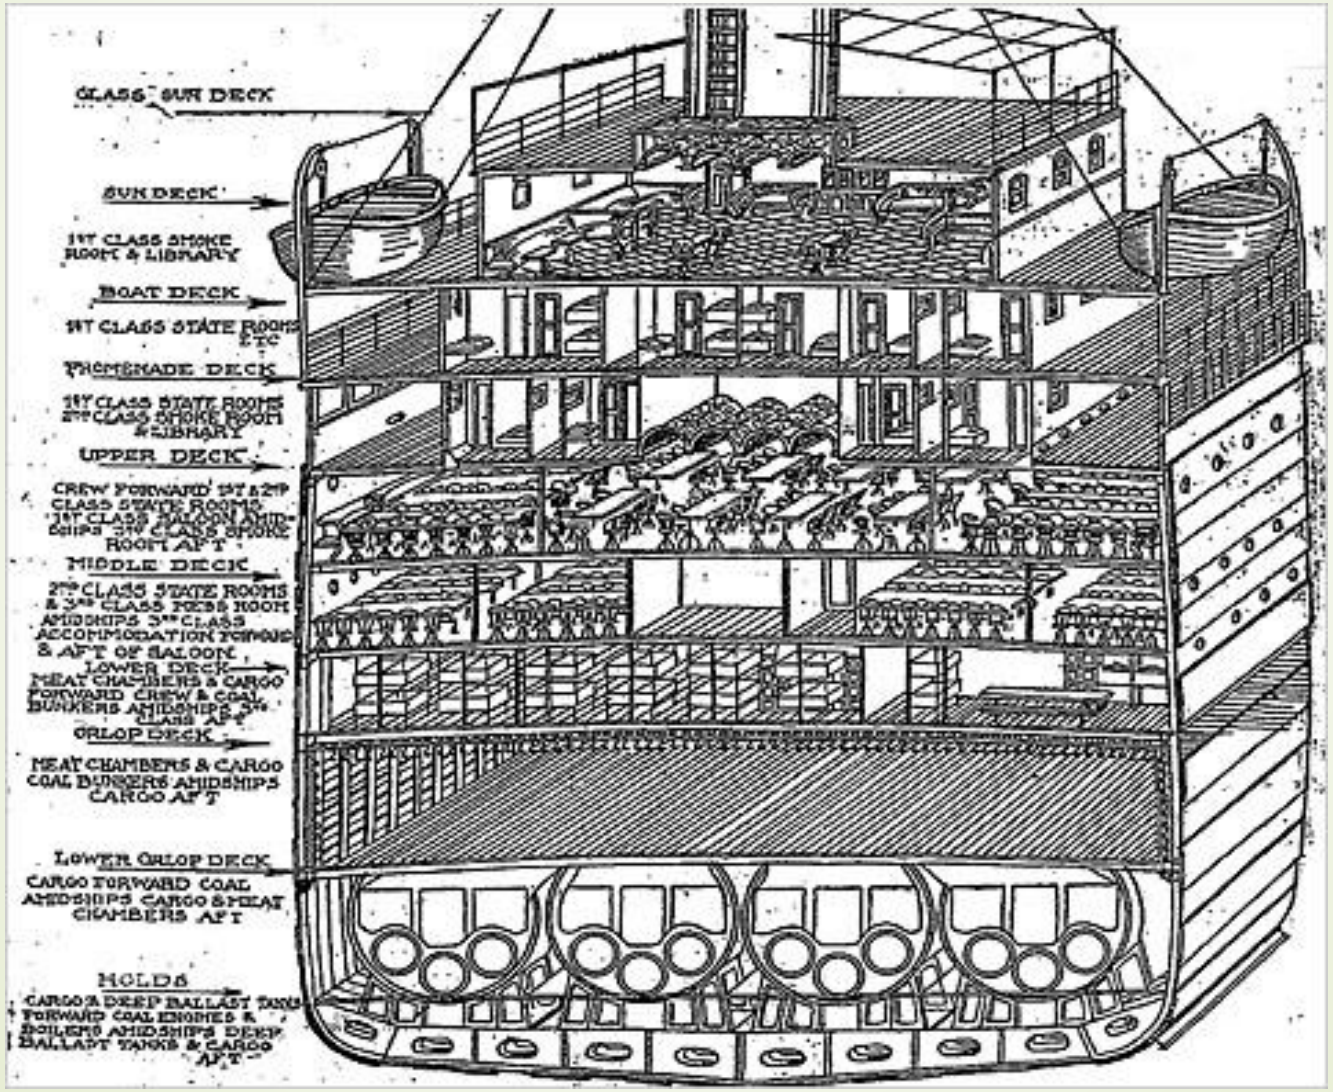

Passenger Survival Rate by Class

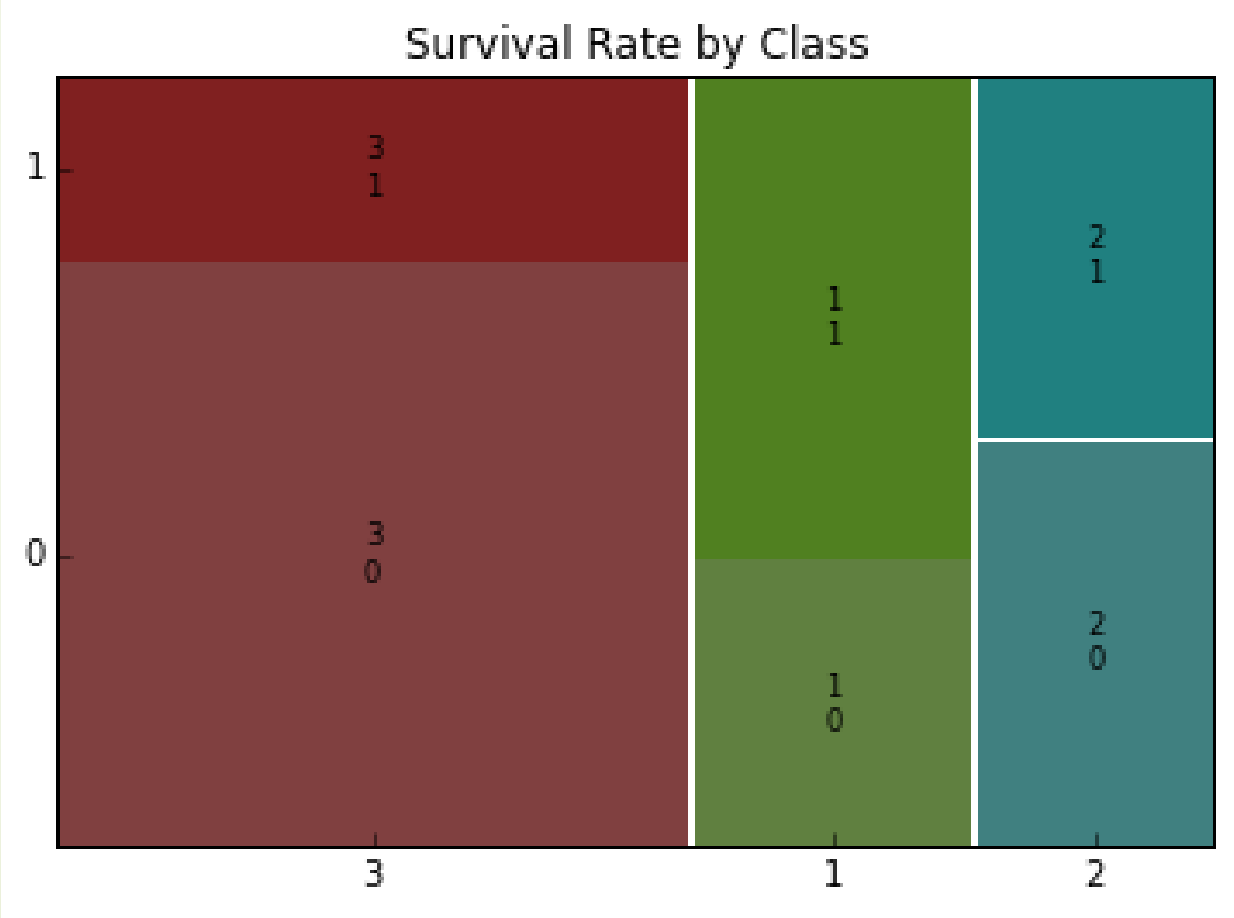

Further Data Manipulation – Port of Embarkation

- The majority of passengers embarked at Southampton in England
- Two variables "Q_Port" and "C_Port" were created to indicate passengers who embarked at Cherbourg in Queenstown

Title Manipulation

- Split titles out from the names column in the data
- Why ?
  - To give more insight into the type of people on board
  - Some highborn people of many different titles
  - Using Title to try and improve average age is slightly more accurate
  - Allows the creation of some additional Categorical Variables

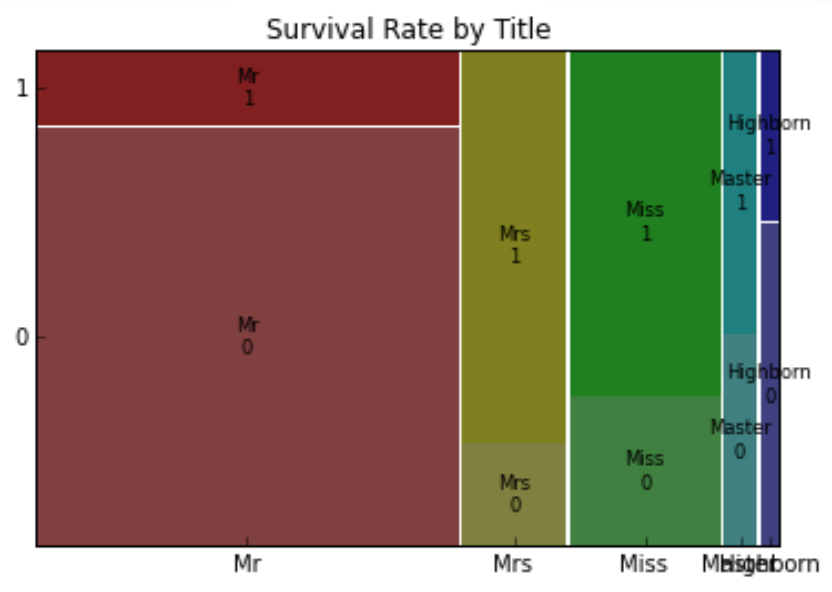

Final Modeling Dataset

	SibSp	Parch	Q_Port	C_Port	Ages	AgeByClass	Family_Size	Elderly	\
0	1	0	0	0	22	66	1	0	
1	1	0	0	1	38	38	1	0	
2	0	0	0	0	26	78	0	0	
3	1	0	0	0	35	35	1	0	
4	0	0	0	0	35	105	0	0	

	Children	First_Class	Second_Class	Mrs	Miss	Highborn	Master	\
0	0	0	0	0	0	0	0	
1	0	1	0	1	0	0	0	
2	0	0	0	0	1	0	0	
3	0	1	0	1	0	0	0	
4	0	0	0	0	0	0	0	

	large_family
0	0
1	0
2	0
3	0
4	0

Model Selection and Methodology

Model Selection

- What type of problem is it?
  - Binary Classification Problem
  - Dead (0) or Alive (1)
- Models Selected:
  - Logistic Regression
  - Random Forests
  - Boosting Trees

Methodology

- Logistic Regression showing poor results after initial testing
- Only optimize the parameters for Random Forest and Boosting Trees
- Used a combination of Recursive Feature Elimination and Grid Search to find the optimal parameters and features to be used in each model
- Random Forests:
  - Optimized for the best number of trees between a range of 10 and 750
  - 5 fold Cross-Validation
- Boosting Trees
  - Tuned the following parameters:
    - Learning Rate Range
    - Sub-Sampling Range
    - Number of Estimators Range

Modeling Results

Logistic Regression

Logit Regression Results

Dep. Variable:	Survived	No. Observations:	891
Model:	Logit	Df Residuals:	879
Method:	MLE	Df Model:	11
Date:	Thu, 09 Apr 2015	Pseudo R-squ.:	0.3584
Time:	09:42:15	Log-Likelihood:	-380.68
converged:	True	LL-Null:	-593.33
		LLR p-value:	2.596e-84

	coef	std err	z	P>|z|	[95.0% Conf.	Int.]
Intercept	-2.6686	0.199	-13.443	0.000	-3.058	-2.280
Master	2.8394	0.525	5.407	0.000	1.810	3.869
Highborn	0.1601	0.487	0.329	0.742	-0.795	1.115
Mrs	3.1331	0.279	11.221	0.000	2.586	3.680
Miss	2.9016	0.258	11.245	0.000	2.396	3.407
C_Port	0.5663	0.239	2.374	0.018	0.099	1.034
Q_Port	0.3454	0.339	1.020	0.308	-0.318	1.009
Elderly	-0.9860	0.461	-2.137	0.033	-1.890	-0.082
Children	-0.0662	0.394	-0.168	0.867	-0.839	0.706
First_Class	2.1606	0.254	8.509	0.000	1.663	2.658
Second_Class	1.1181	0.246	4.545	0.000	0.636	1.600
large_family	-1.6210	0.590	-2.749	0.006	-2.777	-0.465

Random Forests vs Boosting Trees

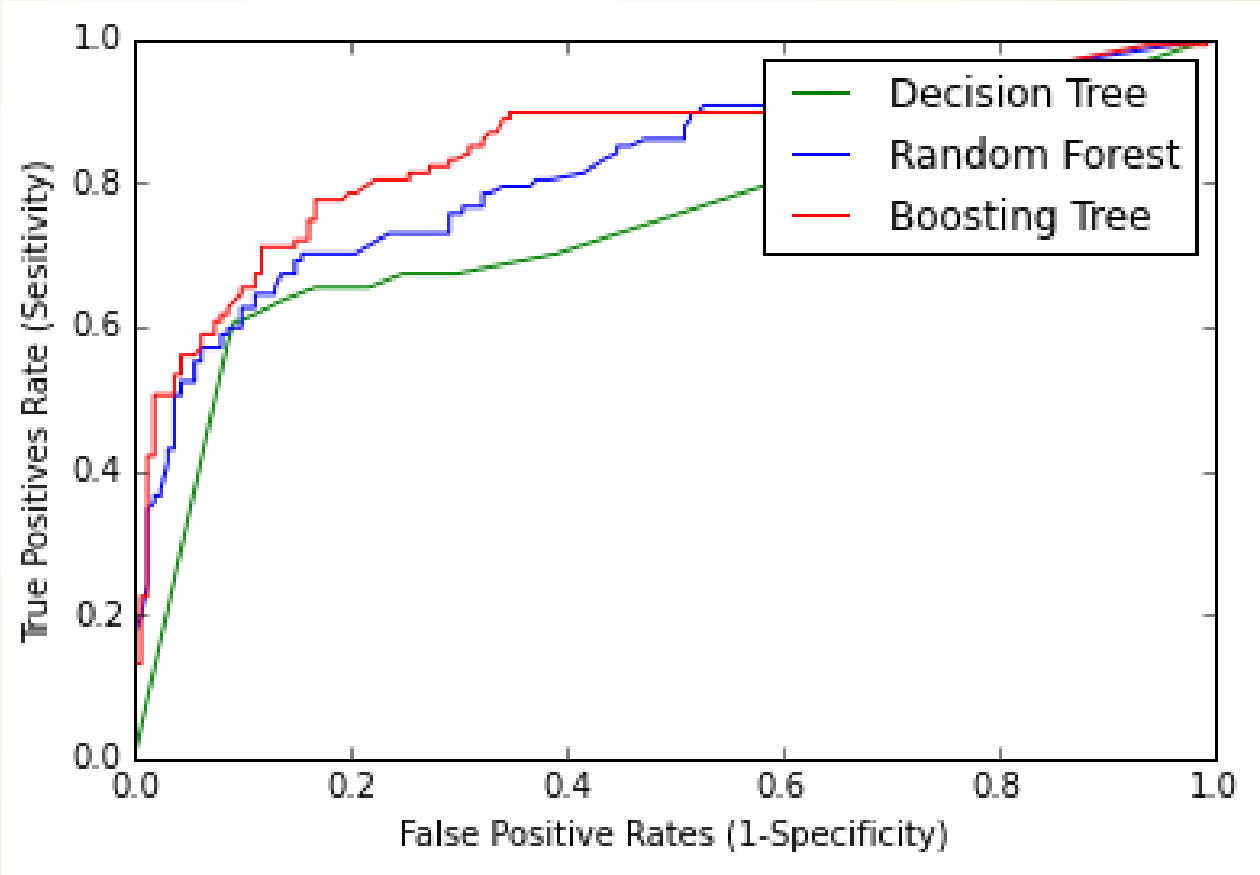

Where I placed on Kaggle?

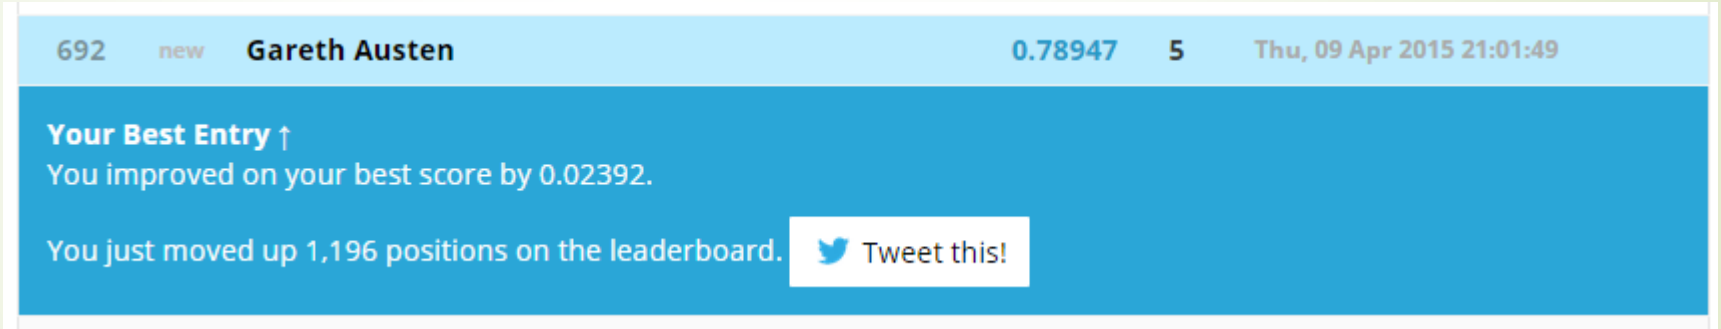

Future Steps

Future Steps and Key Takeaways

- Create further features using combinations of variables
  - Such as combining family size with surname
  - Look at using tickets and cabin numbers
  - Further explore the string variables to try and find other useful features
- Continue to optimize modeling parameters for Random Forests and Boosting Trees
- **Get a faster computer!!**

Questions ?

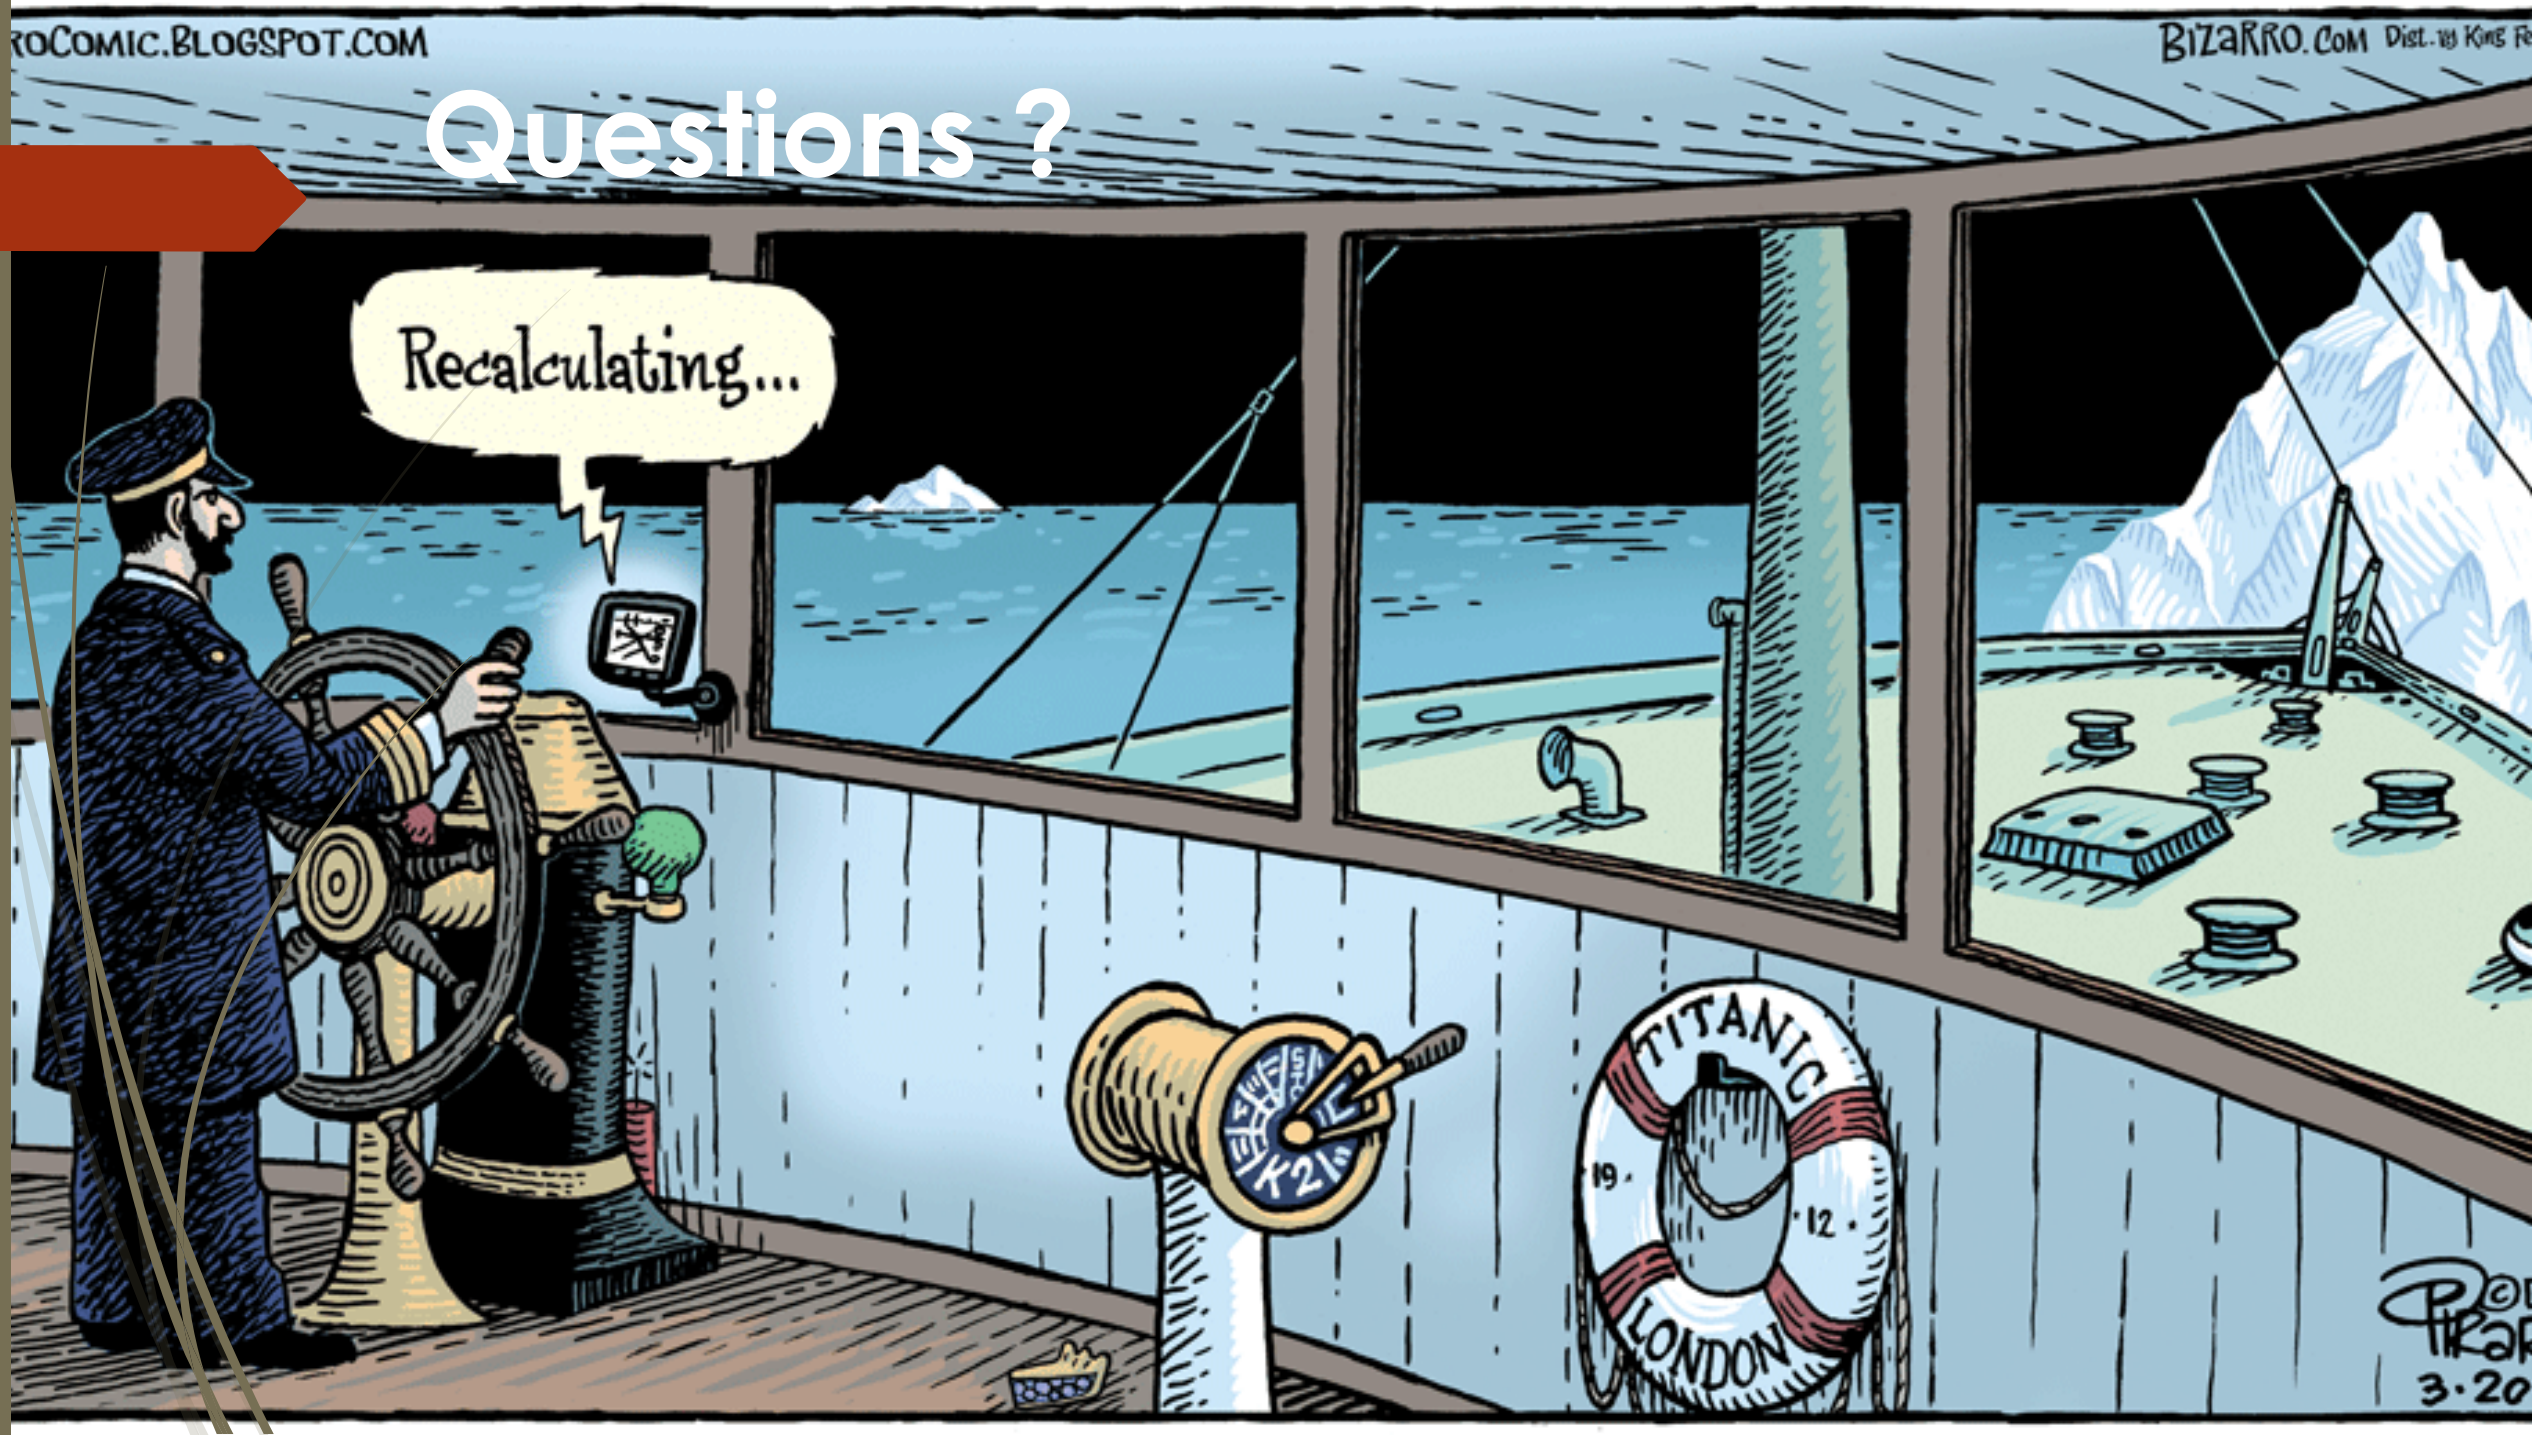
ROCOMIC.BLOGSPOT.COM
BIZARRO.COM
Recalculating...
TITANIC
LONDON
19
12